LAAI - M2 - Homework

Michele Milesi

March 2022

Abstract

Exercise 1

Exercise 1.4

Observe that our model of evaluation allows for combinations whose operators are compound expressions. Use this observation to describe the behaviour of the following procedure:

```
(define (a-plus-abs-b a b)
  ((if (> b 0) + -) a b))
```

Solution The function `a-plus-abs-b` defined above takes in input 2 parameters and returns the sum $a + |b|$. The formal parameters of the fuction are `a` and `b`, while the body of the fuction is composed by a a compound expression. The operator of the expression is a compound expression, indeed, it depends on the value of the parameter `b`, in particular, if the value of b is greater than zero, then it is performed a sum between the operands, while if it is less than or equal to zero, then it is performed a substraction.

When this procedure is called, the formal parameters are substituted by the actual parameters, e.g. when we call the procedure as follows: `(a-plus-abs-b 5 2)` all the instances of `a` in the body of the function are substituted by the value 5 and all the instances of `b` are substituted by 2 and then the body of the function is evaluated. Considering that the operator is a compound expression, the interpreter firts evaluates it. In particular it is a conditional expression, so the interpreter evaluates the predicate `(< b 0)` and if the condition is true, then it will evaluate the *consequent*, otherwise it evaluates the *alternative*. In the previous example, after evaluating the conditional expression the expression become `(+ 5 2)`. At this point the interpreter evaluates the *operator*: since it is a primitive procedure, the interpreter evaluates all the *operands* and it applies the operator to the *arguments* (i.e. the value of the operands). In the example above, it simplifies the expression with `7`. Now the interpreter cannot do any simplification of the expression, indeed, it has to handle a primitive expression, so it does not perform any computation step and it returns the computed value.

Exercise 1.5

Ben Bitdiddle has invented a test to determine whether the interpreter he is faced with is using applicative-order evaluation or normal-order evaluation. He defines the following two procedures:

```
(define (p) (p))

(define (test x y)
  (if (= x 0)
      0
      y))
```

Then he evaluates the expression

```
(test 0 (p))
```

What behaviour will Ben observe with an interpreter that uses applicative-order evaluation? What behaviour will he observe with an interpreter that uses normal-order evaluation? Explain your answer. (Assume that the evaluation rule for the special form `if` is the same whether the interpreter is using normal or applicative order: the predicate expression is evaluated first, and the result determines whether to evaluate the consequent or the alternative expression).

Solution   In applicative-order evaluation, when a procedure is called, the arguments are evaluated first and then applied to the procedure; while in normal-order evaluation the arguments are not evaluated before the procedure call, but they are evaluated when the body of the function is evaluated. In this example, an interpreter which uses applicative-order evaluation first evaluats the arguments, so it evaluates 0 (nothing to do) and `(p)` which causes a loop, indeed, the procedure `p` calls itself and it does not terminate. On the other side, an interpreter which uses normal-order evaluation does not evaluate the arguments before the evaluation of the body of the function, but it passes them without any computation. In this case the procedure call `(test 0 (p))` returns the value 0, indeed, the interpreter evaluates the conditional expression and, since the formal parameter `x` is substituted by the value 0, the condition is true and the consequent (i.e. 0) is returned.

We can observe that the applicative-order evaluation can be useful when the computation of the arguments is heavy and the arguments are used many times in the body of the function, but some arguments can be evaluated even if they are not used and this can cause, as in this particular case, some problems. Instead, normal-order evaluation can be useful when the are used few times or not used, but if they are used many times and their computation is heavy, then the performances are worse than applicative-order evaluation.

Exercise 2

Exercise 1.35

Show that the golden ratio φ is a fixed point of the transformation $x \mapsto 1 + \frac{1}{x}$, and use this fact to compute φ by means of the `fixed-point` procedure.

```
(define tolerance 0.00001)

(define (fixed-point f first-guess)
  (define (close-enough? v1 v2)
    (< (abs (- v1 v2))
       tolerance))
  (define (try guess)
    (let ((next (f guess)))
      (if (close-enough? guess next)
          next
          (try next))))
  (try first-guess))
```

Solution   The *golden ratio* is defined as follows:

$$\varphi = \frac{1 + \sqrt{5}}{2} \approx 1.618$$

It is the fixed point of the transformation $x \mapsto 1 + \frac{1}{x}$ indeed if we apply the transformation to φ we obtain:

$$1 + \frac{1}{\varphi} = 1 + \frac{1}{\left(\dfrac{1 + \sqrt{5}}{2}\right)} = 1 + \frac{2}{1 + \sqrt{5}} = \frac{3 + \sqrt{5}}{1 + \sqrt{5}} = \frac{3 + \sqrt{5}}{1 + \sqrt{5}} \cdot \frac{1 - \sqrt{5}}{1 - \sqrt{5}} = \frac{-2 - 2\sqrt{5}}{-4} = \frac{1 + \sqrt{5}}{2} = \varphi$$

The transformation can be defined in Racket as follows:

```
; definition of transformation x -> 1 + 1/x
(define (transformation x)
  (+ 1 (/ 1 x)))
```

```
; definition of golden ratio
(define phi (/ (+ 1 (sqrt 5)) 2))
```

And the procedure `fixed-point` is used to compute the *golden ratio*, as shown below, we can observe that the procedure `fixed-point` computes φ with a good level of approximation.

```
> (fixed-point transformation 1.1)
1.6180364726455159
> phi
1.618033988749895
```

Exercise 1.36

Modify `fixed-point` so that it prints the sequence of approximations it generates, using the `newline` and `display` primitives shown in Exercise 1.22. Then find a solution to $x^x = 1000$ by finding a fixed point of $x \mapsto \frac{\log(1000)}{\log(x)}$. (Use Scheme's primitive `log` procedure, which computes natural logarithms). Compare the number of steps this takes with and without average damping. (Note that you cannot start `fixed-point` with a guess of 1, as this would cause division by $\log(1) = 0$).

Solution First of all the procedure `fixed-point` has been modified in order to print the sequence of approximations it generates.

```
(define (fixed-point-print-seq f first-guess)
  (define (close-enough? v1 v2)
    (< (abs (- v1 v2))
       tolerance))
  (define (try guess)
    (display guess)
    (newline)
    (let ((next (f guess)))
      (if (close-enough? guess next)
          next
          (try next))))
  (try first-guess))
```

Listing 1: Procedure `fixed-point` which prints the sequence of approximations it generates

Then the transformation $x \mapsto \frac{\log(1000)}{\log(x)}$ has been defined in racket as follows (both with and without average damping):

```
(define (log-transformation x)
  (/ (log 1000) (log x)))
```

Listing 2: Transformation without average-damping

```
; definition of procedure which computes the average
(define (average x y)
  (/ (+ x y) 2))

; definition of transformation x -> log(1000) / log(x)
; with average damping
(define (log-transformation-avg-dmp)
  (fixed-point-print-seq
    (lambda (x)
      (average x (log-transformation x)))
    1.1))
```

Listing 3: Transformation with average damping

After that the two procedures are used to make a comparison between the number of steps, in both cases the initial guess is 1.1 and it can be show how the procedure with average damping takes less time to converge to the solution. Below the approximation of the two methods are shown, the approximation without average damping takes 37 steps, while the approximation with average damping takes 13 steps.

```
; without average damping          ; with average damping
> (fixed-point-print-seq           > (log-transformation-avg-dmp)
   log-transformation 1.1)
1.1                                1.1
72.47657378429035                  36.78828689214517
1.6127318474109593                 19.352175531882512
14.45350138636525                  10.84183367957568
2.5862669415385087                 6.870048352141772
7.269672273367045                  5.227224961967156
3.4822383620848467                 4.701960195159289
5.536500810236703                  4.582196773201124
4.036406406288111                  4.560134229703681
4.95053682041456                   4.5563204194309606
4.318707390180805                  4.555669361784037
4.721778787145103                  4.555558462975639
4.450341068884912                  4.55553957996306
4.626821434106115                  4.555536364911781
4.509360945293209
4.586349500915509
4.535372639594589
4.568901484845316
4.546751100777536
4.561341971741742
4.551712230641226
4.558059671677587
4.55387226495538
4.556633177654167
4.554812144696459
4.556012967736543
4.555220997683307
4.555743265552239
4.555398830243649
4.555625974816275
4.555476175432173
4.555574964557791
4.555509814636753
4.555552779647764
4.555524444961165
4.555543131130589
4.555530807938518
4.555538934848503
```

Exercise 1.37

a. An infinite *continued fraction* is an expression of the form

$$f = \cfrac{N_1}{D_1 + \cfrac{N_2}{D_2 + \cfrac{N_3}{D_3 + \ldots}}}$$

As an example, one can show that the infinite continued fraction expansion with the N_i and the D_i all equal to 1 produces $\dfrac{1}{\varphi}$, where φ is the golden ratio. One way to approximate an

infinite continued fraction is to truncate the expansion after a given number of terms. Such a truncation − a so-called *k-term finite continued fraction* − has the form

$$\cfrac{N_1}{D_1 + \cfrac{N_2}{\ddots + \cfrac{N_k}{D_k}}}$$

Suppose that `n` and `d` are procedures of one argument (the term index i) that return the N_i and D_i of the terms of the continued fraction. Define a procedure `cont-frac` such that evaluating `(cont-frac n d k)` computes the value of the k-term finite continued fraction. Check your procedure by approximating $\frac{1}{\varphi}$ using

```
(cont-frac (lambda (i) 1.0)
           (lambda (i) 1.0)
           k)
```

for succesive values ok `k`. How large must you make `k` in order to get an approximation that is accurate to 4 decimals places?

b. If your `cont-frac` procedure generates a recursive process, write one that generates an iterative process. If it generates an iterative process, write one that generates a recursive process.

Solution

a. The procedure `cont-frac` has been defined with a recursive process. The function has 3 parameters: *(i)* `n` which is the function that returns the element N_i of the continued fraction; *(ii)* `d` which is the function that returns the element D_i of the continued fraction; *(iii)* `k` which is the number of iterations to be performed. In the body of the function is defined the local recursive procedure `cont-frac-rec` which is responsable for computing recursivly the continued fraction. It has two formal parameters: `k` and `counter` which is the counter of the performed iterations; when the `counter` reaches `k` (i.e. the number of iterations to be performed) the base case is reached and the returned value is computed. The procedure `cont-frac` calls the procedure `cont-frac-rec` with initial parameters `k` and 0, in this way we are sure that the number of iterations performed will be `k` (with `counter` from 1 to `k`). The code is shown below:

```
(define (cont-frac n d k)
  (define (cont-frac-rec k counter)
    (if (= counter k)
        (/ (n counter) (d counter))
        (/ (n counter)
           (+ (d counter)
              (cont-frac-rec k (+ counter 1))))))
  (cont-frac-rec k 1))
```

Listing 4: Recursive procedure `cont-frac`

Then it has been tested by approximating $1/\varphi \sim 0.6180$. When $k = 11$ or grater, the approximation is accurate to 4 decimals places.

b. Since before the procedure generates a recursive process, it has been rewritten in order to generate an iterative process. The procedure defines the local procedure `iter` which takes in input the number of iterations to be performed (`k`) and partial result computed so far. The idea is to begin from the last fraction (i.e. N_k/D_k) and then proceding backward to compute all the other fractions. In particular, the idea is to compute the quantity $D_{i-1} + \frac{N_i}{Q_i}$ where Q_i is the quantity computed so far. Indeed, the procedure `iter` is called with initial parameters `k` and $Q_k = D_k$, after the first iteration, the computed value $Q_{k-1} = D_{k-1} + \frac{N_k}{D_k}$. In the last step (i.e. $k = 1$) the procedure returns the ration N_1/Q_1 that is exaclty the continued fraction to be computed.

```
(define (cont-frac-iter n d k)
  (define (iter k res)
    (if (= k 1)
        (/ (n k) res)
        (iter (- k 1)
              (+ (d (- k 1))
                 (/ (n k) res)))))
  (iter k (d k)))
```

Exercise 1.38

In 1737, the Swiss mathematician Leonhard Euler published a memoir *De Fractionibus Continuis*, which included a continued fraction expansion for $e-2$, where e is the base of the natural logarithms. In this fraction, the N_i are all 1, and the D_i are successively $1, 2, 1, 1, 4, 1, 1, 6, 1, 1, 8, \ldots$. Write a program that uses your `cont-frac` procedure from Exercise 1.37 to approximate e, based on Euler's expansion.

Solution   The two procedures `cont-frac` and `cont-frac-iter` were be used to approximate e in order to verify that the results matched. The procedure `euler-number` uses the procedure `cont-frac` (which generates a recursive process), while the procedure `euler-number-iter` uses the procedure `cont-frac-iter`. To get the values of N_k and D_k, two procedures have been defined: `get-n` and `get-d`. The first one is trivial and returns 1 at each iteration step, while the latter is more complex. To get the element of the sequence $1, 2, 1, 1, 4, 1, 1, 6, 1, 1, 8, \ldots$ it is possible to use the function $seq : \mathbb{N}^+ \to \mathbb{N}^+$ described below:

$$seq(i) = \begin{cases} i & \text{if } 1 \leq i \leq 2 \\ \left(\frac{i-2}{3} + 1\right) \cdot 2 & \text{if } i - 2 \bmod 3 = 0 \\ 1 & \text{otherwise} \end{cases}$$

So we defined the procedures `get-n` and `get-d` as follows:

```
(define (get-n i) 1)
(define (get-d i)
  (if (= (remainder (- i 2) 3) 0)
      (* (+ (quotient
             (- i 2)
             3)
           1)
         2)
      (if (<= i 2) i 1)))
```
Listing 5: Definition of the procedures `get-n` and `get-d`

Since the sequence of D_i is used to approximate $e-2$, we can approximate e by approximating $e-2$ and than by adding 2 to the computed quantity. This is the way the procedures `euler-number` and `euler-number-iter` approximate e.

```
(define (euler-number k)
  (+ (cont-frac get-n get-d k) 2))

(define (euler-number-iter k)
  (+ (cont-frac-iter get-n get-d k) 2))
```

Both the procedures return the value $23225/8544 \sim 2.7182 \approx e$.

Exercise 3

Exercise 3.2

Explain why (in terms of the evaluation process) these two programs give different answers (i.e. have different distributions on return values):

```
(define foo (flip))
(list foo foo foo)
```

```
(define (foo) (flip))
(list (foo) (foo) (foo))
```

Solution   In the first program we are defining `foo` as a variable and we are assigning it the value of the evaluation of the expression `(flip)`, indeed, the value of `foo` is either `#t` or `#f`. After that we create a list which contains three time the value of `foo`, so if `foo` has value `#t`, then we are defining a the list: `'(#t #t #t)`; otherwise we are creating the list: `'(#f #f #f)`.

In the second program, we are defining a new procedure `foo` which is a wrapper for the procedure `flip`. This means that each time the procedure `foo` is called, also the procedure `flip` is called. For this reason, when we define the list, the expression `flip` is evaluated three times and, since it is a non-deterministic procedure, the three elements of the list can be different.

Exercise 3.5

Here is a modified version of the tug of war game. Instead of drawing strength from the continuous Gaussian distribution, strength is either 5 or 10 with equal probability. Also the probability of laziness is changed from 1/4 to 1/3. Here are four expressions you could evaluate using this modified model:

```
(define strength (mem (lambda (person) (if (flip) 5 10))))
(define lazy (lambda (person) (flip (/ 1 3))))

(define (total-pulling team)
  (sum
   (map (lambda (person) (if (lazy person)
                             (/ (strength person) 2)
                             (strength person)))
        team)))

(define (winner team1 team2)
  (if (< (total-pulling team1) (total-pulling team2))
      team2
      team1))

(winner '(alice) '(bob))                              ;; expression 1

(equal? '(alice) (winner '(alice) '(bob)))       ;; expression 2

(and (equal? '(alice) (winner '(alice) '(bob)))  ;; expression 3
     (equal? '(alice) (winner '(alice) '(fred))))

(and (equal? '(alice) (winner '(alice) '(bob)))  ;; expression 4
     (equal? '(jane) (winner '(jane) '(fred))))
```

a. Write down the sequence of expression evaluations and random choices that will be made in evaluating each expression.

b. Directly compute the probability for each possible return value from each expression.

c. Why are the probabilities different for the last two? Explain both in terms of the probability calculations you did and in terms of the "causal" process of evaluating and making random choices.

Solution

a. First of all, the procedures are defined and the interpreter associates the name of the procedures with their definition in the *global environment*. Then the interpreter evaluates the four expressions in the order in which they are written:

Expression 1 `(winner '(alice) '(bob))`

The first step is to retrieve the body of the procedure `winner` and then the formal parameters are substituted by the actual parameters `'(alice)` and `'(bob)`. Now the interpreter has to evaluate the following expression:

```
(if (< (total-pulling '(alice)) (total-pulling '(bob)))
    '(bob)
    '(alice)))
```

Listing 6: Body of procedure `winner` with actual parameters

The following step is to evaluate the conditional expression, so it starts by evaluating the *predicate*: `(< (total-pulling '(alice)) (total-pulling '(bob)))`. The interpreter has to deal with the primitive predicate `<`, so it needs to evaluate the arguments and than apply the predicate to the evaluated arguments. The first argument to evaluate is `(total-pulling '(alice))`: the body of the procedure `total-pulling` is retrieved in the *global environment* and it is substituted by the expression, then the actual parameter is applied to the body of the procedure. The expression `(total-pulling '(alice))` becomes as follows:

```
(sum
 (map (lambda (person) (if (lazy person)
                           (/ (strength person) 2)
                           (strength person)))
      '(alice)))
```

Listing 7: Body of procedure `total-pulling` with actual parameters

The procedure `sum` takes as argument a list of numbers and returns the sum over the elements of the list. In this case the argument of the procedure is another expression, so the interpreter has to evaluate the argument and then apply it to the procedure `sum`. Now the expression to be evaluated is the following:

```
(map (lambda (person) (if (lazy person)
                          (/ (strength person) 2)
                          (strength person)))
     '(alice))
```

Listing 8: Argument of procedure `sum` in Listing 7

The procedure `map` takes as arguments a *lambda* expression and a list. The procedure applies the *lambda* expression to each element of the list and returns the list of the results of the *lambda* expression on the elements of the list in input. So the first step is to evaluate the *lambda* expression with argument `'alice`. The resulting expression is the body of the *lambda* expression with actual parameter `'alice`:

```
(if (lazy 'alice) (/ (strength 'alice) 2) (strength 'alice))
```

Listing 9: Application of `map` prcedure of Listing 8

Since the interpreter has to evaluate a conditional expression, the interpreter firsts evaluate the *predicate*. To evaluate the call `(lazy 'alice)`, the interpreter first takes the body of the procedure and then substitutes all the occurences of formal parameter by the value of the actual parameter. The resulting expression is:

```
(flip (/ 1 3))
```

Listing 10: Body of procedure `lazy` when evaluating Listing 9

The procedure `flip` is a probabilistic procedure, it can take as parameter the probability to return `#t`. In our case, the interpreter has to evaluate the expression `(/ 1 3)` and than apply the result to the procedure `flip`. We suppose that the result of the call `(flip 1/3)` is `#f`, so the result of the evaluation of the *predicate* in the conditional expression is `#f`. Since the *predicate* is false, the interpreter has to evaluate the *alternative* expression `(strength 'alice)`. It is a procedure call, so the interpreter has to retrieve the body of the procedure and then to substitute the formal parameter by the actual one. The result is that the interpreter now has to evaluate the following expression:

```
(if (flip) 5 10)
```
Listing 11: Body of procedure `strength` when evaluating Listing 9

The procedure `strength` is a particular type of procedure, indeed, it is memoized through the procedure `mem`. This means that the first time the procedure is called, the result of the call is memorized and each time the procedure is called with same the same argument, the interpreter does not evaluate the function, but it retrieves the already computed result in the *global environment*. This is the first time that `(strength 'alice)` is evaluated, so the procedure has to be evaluated in a "standard" way. The body of the procedure is composed by a conditional expression, the procedure `flip` is evaluated and we suppose that the result is `#t`, so the procedure `strength` returns the value 5.

Thus, the expression `(strength 'alice)` is evaluated with 5; since the second argument of the procedure `map` is a list with only one argument (i.e. `'alice`) the result is a list with only one element (i.e. 5). This list is the argument of the procedure `sum` and the interpreter has to evaluate the following expression:

```
(sum '(5))
```
Listing 12: Expression to be evaluated after evaluating the expression in Listing 8

Which is evaluated with 5. This is the result of the evaluation of the call (`total-pulling '(alice)`), now the interpreter has to evaluate the second argument of the expression (`< 5 (total-pulling '(bob))`), since 5 is a primitive expression and it cannot be reduced anymore. The evaluation of the expression (`total-pulling '(bob)`) is the same as the expression (`total-pulling '(alice)`), but, since there are some probabilistic cases, the result can be different. In particular, we suppose that the evaluation of (`lazy 'bob`) returns the value `#t`, so the interpreter has to evaluate the expression (`/ (strength 'bob) 2`) instead of (`strength 'bob`). The interpreter has to evaluate a primitive procedure, so it has to evaluate first the operands and then to apply the operator to them. We suppose that the evaluation of (`strength 'bob`) returns the value 5. It is important to notice that it is the first time that the procedure `strength` is called with actual parameter `'bob`, so the evaluation is "standard" and the result of the call is memorized in the *global environment*. So the expression (`/ (strength 'bob) 2`) is evaluated with 2.5.

Now the situation is represented in the following way:

```
(if (< 5 2.5)
    '(bob)
    '(alice)))
```
Listing 13: Expression to be evaluated after evaluating (`total-pulling 'bob`)

Since $5 > 2.5$, the *predicate* is false and the value `'(alice)` is returned. Thus the *expression 1* is evaluated with the value `'(alice)`.

Expression 2   `(equal?  '(alice) (winner '(alice) '(bob)))`

Expression 3

b.

c.

Exercise 3.6

Use the rules of probability, described above, to compute the probability that the geometric distribution defined by the following stochastic recursion returns the number 5.

```
(define (geometric p)
  (if (flip p)
      0
      (+ 1 (geometric p))))
```

Solution   The procedure computes the number of consecutive *false* (`#f`) results. Since each coin toss (i.e. `flip`) is independent, the probability of getting five consecutive *false* results (and the sixth one *true*) is given by:

$$P(geometric = 5) = (1 - p)^5 \cdot p$$

The formula comes from the fact that the probability of getting a *true* value from the procedure `flip` is given by p, so the probability of getting a *false* value from `flip` is $(1-p)$. So the procedure `geometric` computes the number of trials needed to get the first occurence of success (i.e. `#t`). Each trial has the same probability of success p. For this reason the computed probability is equivalent to the geometric distribution with success probability p and with the first occurence of success at the sixth trial.

To check the formula, some samples have been generated in order to approximate the probability to get five consecutive *false* results. The experiment consists of generating 100000 samples with probability $P(true) = P(false) = 0.5$. The following procedures are defined in order to implement the experiment: *(i)* `model` is a wrapper for the procedure `geometric`; *(ii)* `count-5` takes in input the list of samples and returns the number of occurences which have value 5. Then the samples are generated and the statistics are computed.

```
; number of samples we want to generate
(define n-samples 100000)

; model used to generate the samples
(define (model)
  (define p 0.5)
  (geometric p))

; procedure which counts the number of samples with value 5
(define (count-5 l)
  (if (null? l)
      0
      (if (= (car l) 5)
          (+ 1 (count-5 (cdr l)))
          (count-5 (cdr l)))))

; sampling
(define experiment (repeat model n-samples))

; ratio between the number of samples with value 5
; and the total number of samples
(/ (count-5 experiment) n-samples)

; histogram of the results
(hist experiment)
```

Listing 14: Experiment to approximate the probability of getting the first occurence of success at the sixth trial

The probability computed by hand is $P(geometric = 5) = 0.5^5 \cdot 0.5 = 0.5^6 = 0.015625$, while the probability computed by the program is $P_{program}(geometric) = 1569/100000 = 0.01569$. The two probabilities are very similar, so we can conclude that the calculation of the probability is correct. The histogram of the generated samples is shown in Figure 1.

Exercise 3.7

Convert Table 1 to a compact Racket program.

Figure 1: Histogram of geometric experiment: the *x-axis* represents the values generated during the sampling phase; the *y-axis* represents the number of samples which have a specific value.

A	B	P(A, B)
F	F	0.14
F	T	0.06
T	F	0.4
T	T	0.4

Table 1: Probabilities to be computed with a Racket program

Hint: fix the probability of A and then define the probability of B to depend on whether A is true or not. Run your Church program and build a histogram to check that you get the correct distribution.

```
(define a ...)
(define b ...)
(list a b)
```

Solution The `a-b-model` has been defined as follows:

```
(define (a-b-model)
  (define a
    (flip 0.8))
  (define b
    (if a
        (flip 0.5)
        (flip 0.3)))
  (list a b))
```

Listing 15: Model to compute the probabilities of A and B

The model does not contain the `rejection-sampler` because we do not need to compute a conditional probability. The `a-b-model` defines first the variable `a` which has probability 0.8 to be *true*: this probability can be computed by adding the last two rows of the Table 1, indeed, the value of A in the first two rows is *false*, while in the last two is *true*. Then the probability of the variable

11

b depends on the value of the variable `a`, indeed, if `A` is *false*, then the probability that `B` is *true* $\frac{0.06}{0.06+0.14} = 0.3$; while if `A` is *true*, then la probability that `B` is true is 0.5.


Figure 2: Histogram of A-B experiment: the *x-axis* represents the values generated during the sampling phase; the *y-axis* represents the number of samples which have a specific value.

The experiment consists of generating 100000 samples, the histogram of generated samples is shown in Figure 2. It is possible to observe that the samples are distributed as the probability distribution defined in Table 1, indeed, the number of samples with value false both for the variable `A` and `B` is about 14000; the number of samples with `A` *true* is about 6000 and the number of samples with value *(true, false)* and *(true, true)* is about 40000 each. The number of samples is not exact because we are approximating the probability distribution by sampling.

Exercise 4

Exercise 4.1

What are (`bernoulli-dist p`), (`normal-dist` μ σ) exactly? Are they real numbers (produced in a random way)? We have seen that `flip` is a procedure with a probabilistic behaviour. Is, e.g., (`normal-dist` μ σ) something similar? Try to evaluate (`normal-dist 0 1`)

Solution (`bernoulli-dist p`) and (`normal-dist` μ σ) are two structures which represent two different distribution, the first one represents the *bernoulli* distribution with success probability p; instead, the latter object represents the *normal* distribution with parameters μ and σ as mean and variance. They are not real numbers, they are structures which can be used to generate numbers according to the distribution they represent. To generate a sample from a distribution object it necessary to use the procedure `sample` which takes as argument a distribution object and returns the generated value.

The procedure `flip` has a probabilistic behaviour, but it is different from the call (`normal-dist 0 1`), indeed, the first one is a procedure which returns the value `#t` with probability 0.5, while the second one is a structure, so the evaluation is different. Since it is a structure, so when the call is done, the interpreter returns a distribution object which has as parameters the arguments of the call, e.g. in our case it returns an object which represents a normal distribution with $\mu = 0$ and $\sigma = 1$.

The procedure `flip` has a similar behaviour to the call (`sample (bernoulli-dist 0.5)`) that is equivalent to the call (`bernoulli 0.5`). One difference is that `flip` returns a value which is

either **#f** or **#t**, while (**bernoulli 0.5**) returns a value which is either 0 or 1.

Exercise 4.2

Evaluate

```
(dist? (normal-dist 0 1))

(dist? (bernoulli-dist 0.5))

(dist? flip)
```

What is the difference between flip and (bernoulli-dist 0.5)?

Solution

1. Evaluation of (**dist?** (**normal-dist 0 1**)): The interpreter first evaluate the procedure name **dist?**, then it evaluates it arguments. The only argument the interpreter has to evaluate is (**normal-dist 0 1**). To evaluate this argument, it evaluates first element of the list and than it evaluates the arguments of expression. The expression (**normal-dist 0 1**) is a particular expression, indeed, it is a struct, so the interpreter creates a new instance of the struct **normal-dist** with parameters 0 and 1. The returned instance is the actual parameter of the procedure **dist?** which returns **#t** if the argument is a distribution object, **#f** otherwise. In this particular case the returned object is a distribution object, so the evaluation of the expression is **#t**.

2. Evaluation of (**dist?** (**bernoulli-dist 0.5**)): The evaluation of this expression is very similar to the previous one, indeed, the procedure **dist?** has the same behaviour as before and the evaluation of (**bernoulli-dist 0.5**) is similar to the evaluation of (**normal-dist 0 1**). In both cases the interpreter has to deal with a struct, so it evaluates the *constructor* and returns an instance of the structure type. In this case it returns a distribution object which represents a bernoulli distribution with success probability of 0.5. Also in this case the final evaluation is **#t**.

3. Evaluation of (**dist?** **filp**): In this case the result is different, indeed, the procedure **flip** is not a distribution object, but it is a procedure with probabilistic behaviour. For this reason, when the interpreter evaluates (**dist?** **flip**), it returns **#f**.

The difference between **flip** and (bernoulli-dist 0.5) is that the first one is a procedure that can be called and its evaluation can return the value **#f** or **#t** both with probability 0.5. Instead, the second expression is the call to a constructor of the structure **bernoulli-dist** and the evaluation returns a distribution object of the structure type, i.e. it returns an instance with parameter 0.5. From this object it is possible to return some samples by the procedure according to the bernoulli distribution with success probability equal to 0.5.

Exercise 5

Exercise 5.4

Probabilistic Models of Cognition - Exercise 4

Solution

Exercise 6

Exercise 6.1

To see the problems of rejection sampling, consider the following variation of the previous example:

```
(define baserate 0.1)
(define (take-sample)
  (rejection-sampler
   (define A (if (flip baserate) 1 0))
```

```
    (define B (if (flip baserate) 1 0))
    (define C (if (flip baserate) 1 0))
    (define D (+ A B C))
    (observe/fail (>= D 2))
    A))
```

Try to see what happens when you lower the basesate. What happens if we set it to 0.01? And to 0.001?

Solution In order to assess the differences between different baserates, it has been developed a program which generates 100 samples for each baserate and computes the histogram of the results. The code is shown below:

```
(define (model baserate)
  (define (take-sample)
    (rejection-sampler
     (define A (if (flip baserate) 1 0))
     (define B (if (flip baserate) 1 0))
     (define C (if (flip baserate) 1 0))
     (define D (+ A B C))
     (observe/fail (>= D 2))
     A))
  (take-sample))

; experiment with baserate = 0.1
(hist (repeat (model 0.1) 1000))

; experiment with baserate = 0.01
(hist (repeat (model 0.01) 1000))

; experiment with baserate = 0.001
(hist (repeat (model 0.001) 1000))
```

The procedure `model` has been developed in order to be able to pass the baserate as parameter of the procedure. So the procedure `model` is a wrapper for the procedure `take-sample`. In this way it is possible to call the procedure `take-sample` with different baserates by passsing a different parameter to the procedure `model`. Then three different experiments are run: *(i)* The baserate is set to 0.1; *(ii)* The baserate is set to 0.01; *(iii)* The baserate is set to 0.001. By observing the histograms of the results of the different experiments we can conclude that the computed probability is more or less the same in all three cases, but the main difference is that the time of execution is completely different. In particular the first example (i.e. $baserate = 0.1$) is faster than the other two cases. Furthermore the case with $baserate = 0.01$ takes much less time that the third case.

This behaviour is due to the probability to get at least two successful results, in particular the lower is the baserate the lower is the probability that the procedure `flip` returns `#t` so the lower is the probability that `A`, `B` and `C` are equal to one. Since we are approximating the posterior probability $P(A|D >= 2)$ by rejection sampling, all the samples which do not agree with the evidence (i.e. $D = 2$) are discarded. When the baserate is high the probability of getting $D >= 2$ increases, so it is less likely that the sample is discarded, instead if the baserate is low, then the probability to discard a sample increases because it is more probable that the variables `A`, `B` and `C` are equal to zero.

The results of the approximated probabilities are shown in Figure 3, Figure 4 and Figure 5. It is possible to observe that $P(A = 0|D >= 2) \approx 1/3$ and $P(A = 0|D >= 2) \approx 2/3$.

Exercise 7

Exercise 7.1

Proposition. *The functional and matrix-based definitions of a MC are equivalent.*

Proof sketch. Given $c : X \to D(X)$, with $X = \{x_1, ..., x_n\}$, we construct the matrix P_c as $P_c(i, j) = c(x_i)(x_j)$. Vice versa, given P, we define $c_P(x_i)(x_j) = P(i, j)$. $\square$

Figure 3: Approximate posterior probability $P(A|D >= 2)$ with $baserate = 0.1$

Figure 4: Approximate posterior probability $P(A|D >= 2)$ with $baserate = 0.01$

Figure 5: Approximate posterior probability $P(A|D >= 2)$ with $baserate = 0.001$

Instructions *Complete the above proof. Prove, in particular that for any $x \in X$, $c_P(x)$ is indeed a distribution; that c_P is a stochastic matrix; and that $P_{cP} = P$ and $c_{Pc} = c$.*

Solution

Exercise 7.2

Prove that $c(x) = c^*(\delta_x)$.

Solution

Exercise 7.3

Prove that $c^*(\psi) = \psi(P_c)$.

Solution

Exercise 7.4

Prove that $f \psi$ satisfied DBC, then ψ is stationary for P.

Solution